



Socket Programming

Programmazione di Reti

Ing. Chiara Contoli, PhD

chiara.contoli@unibo.it

*Corso di Laurea Triennale in
Ingegneria e Scienze Informatiche*

Outline

Socket Programming:

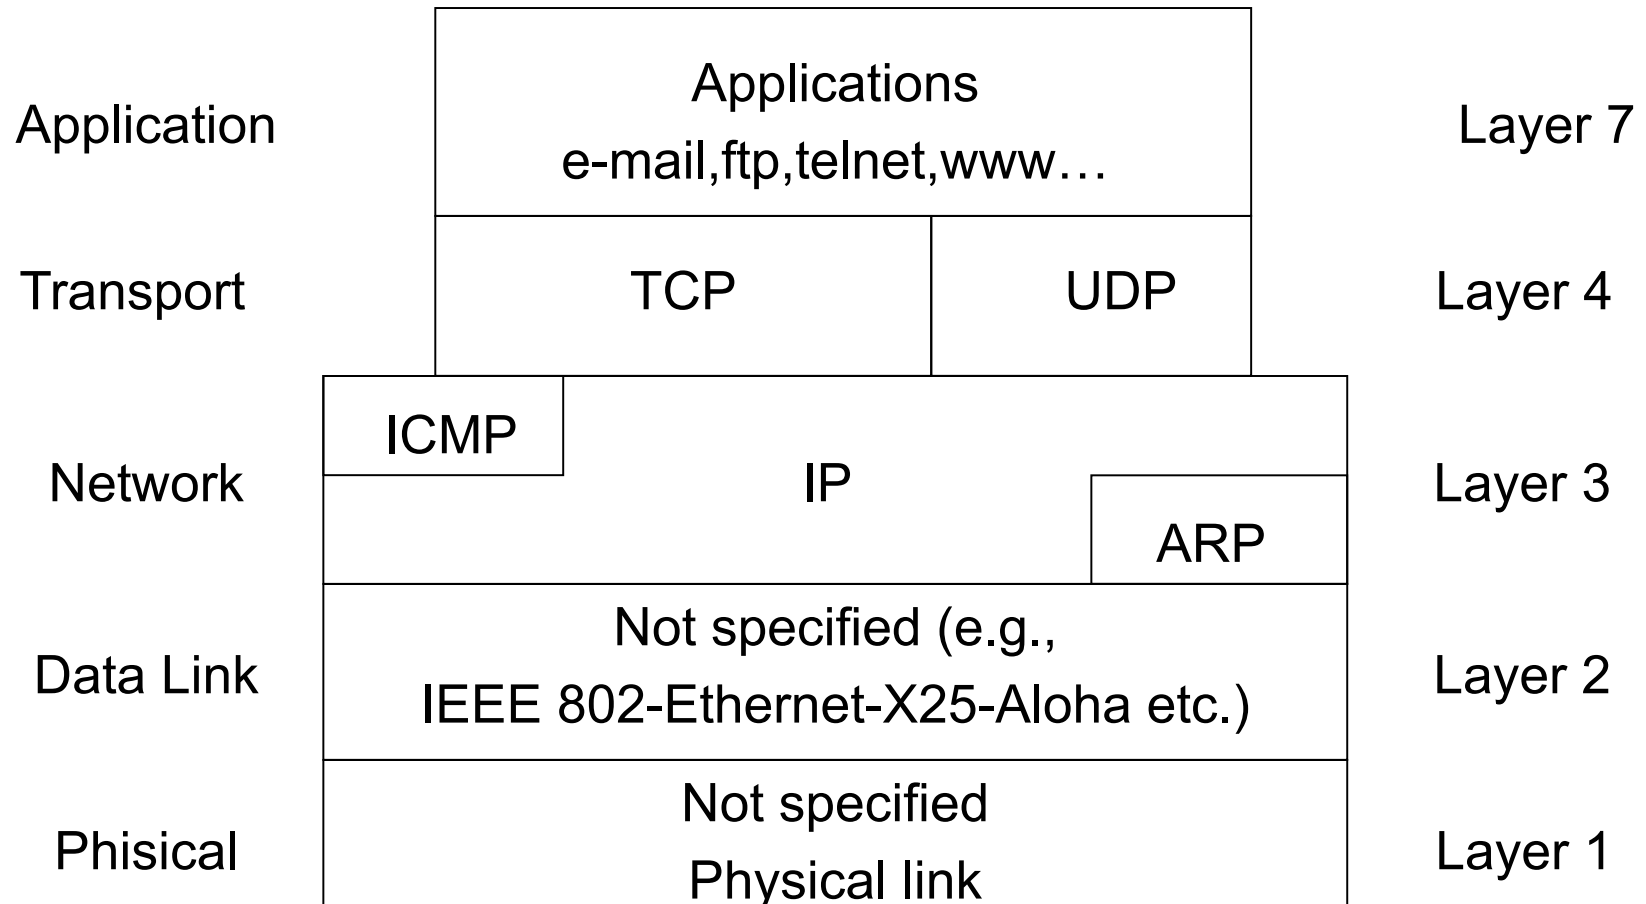
- Basic concepts
- UDP socket
- Client/Server interaction

C language Socket Programming

- Simple UDP Client application
- Simple UDP Server Application

ISO-OSI

Distributed systems: communications between different processes via messages exchange



IP Network

End-to-end communication: IP network allows (distributed) hosts located at the network edges to communicate between each others


Processes use layer 7 protocol in order to exchange information

Communication occurs by exchanging packets between remote hosts

- Layer 4 protocol adoption

Network routes packets via layer 3

Packets are directly delivered to layer 2



Socket programming

Goal: learn how to build client/server applications that communicate via socket

Socket API

- Introduced in 1981 in UNIX BSD 4.1
 - We will mainly take as reference the Linux kernel networking stack
- **Client / server** paradigm
- Sockets are explicitly created and utilized from the application
- Abstract pretty similar to file access
- Two types of transport via socket API:
 - **UDP** (unreliable)
 - **TCP** (reliable, byte stream)


Socket Programming: introduction

Today is pretty likely that systems use some form of networking.

Layer 2, 3 and 4 are all handled by the Linux *kernel*. It handles both *incoming* and *outgoing* packets. Incoming packets might be delivered *locally* or *forwarded*.

Application layer is not handled by the kernel, it is handled by the *userspace*.

This implies that (at least) two interfaces are required: one towards the userspace (i.e., towards applications) and one towards the kernelspace (i.e., towards the network).



Such interface is provided by the so called *Socket API*.

Socket Programming: basic concepts

Socket: a *local host* “interface”, *created by applications (controlled by operating system, OS)*, through which the process of an application can **send** and **receive** messages to / from the process of another application

The **client** must contact the **server**

Both client and server need to have their own sockets through which they can send and receive datagrams

To each socket is associated a **local port number**

Client **needs to know** the server **IP address** and the **port number** on which the application (and related socket) are listening to

Socket Programming: basic concepts

Several types of sockets exist:

- Stream sockets: byte-stream oriented and reliable communication
- Datagram sockets: message oriented, unreliable communication
- Raw sockets: allow direct access to IP layer (packets are sent/received directly to / from IP layer)
- ...

Sockets are defined in a *communication domain*, which identifies:

- Address identification (i.e., socket address format)
- Communication range (i.e., same host or distributed hosts)

Several domains exists:

- Unix domain (i.e., communication occurs between application running on the same host)
- Internet domain (i.e., communication occurs among applications running on different hosts) → TCP/IP protocol stack

Socket Programming: basic concepts

Socket API provides several methods, among which:

- *Socket()*: creates a new socket
- *Bind()*: associates a socket with a local IP address
- *Listen()*: allow to receive connection from other socket (not used in datagram sockets)
- *Accept()*: accepts a connection from a peer socket (used in connection oriented socket)
- *Connect()*: allow to establish a connection to a peer socket (used in connection oriented socket)

Specific system call are also provided for socket I/O:

- *Send()*: send byte-streams
- *Recv()*: receive byte-streams
- *Sendto()*: send messages
- *Recvfrom()*: receive messages

C Socket

Sockets in C language are implemented through a standard service interface to be invoked on the Operating System

- Socket: communication endpoint
- Socket descriptor: equivalent to a file descriptor in UNIX

Communication domain defines communication semantics

- Each communication domain may support different protocols
 - *UNIX (AF_UNIX)*: local communication via pipe
 - *Internet (AF_INET)*: remote communication via internet protocols (TCP/IP)
 - *XROX (AF_NS)*: Xerox protocol
 - *AF_CCITT (AF_CCITT)*: protocollo X.25

Socket can have different communication domains

Socket Libraries

Several libraries are needed: `<sys/socket.h>`,
`<arpa/inet.h>`, `<netinet/in.h>`

From `<sys/socket.h>`

- `int socket(int domain, int type, int protocol);`
 - Returns (socket) file descriptor on success, -1 on error
 - Parameters: i) *domain*: communication domain, ii) *type*: socket type, iii) *protocol*: protocol to be used
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Returns 0 on success, -1 on error
 - Parameters: i) *sockfd*: socket file descriptor, ii) *addr*: address to be bound to the socket; *addrlen*: size of the ii)

Internet domain socket address

Two types: IPv4 and IPv6. Our focus is on IPv4; address is stored in a data structure: `sockaddr_in` (struct `sockaddr_in`, defined in `<netinet/in.h>`)

```
struct sockaddr_in { // IPv4 socket address
    sa_family_t sin_family; // Address family
    in_port_t sin_port; // Port
    struct in_addr sin_addr; // IPv4 address
    ....
}
```

```
struct in_addr { // IPv4 address (4-byte)
    in_addr_t s_addr;
}
```



UDP Socket

Programmazione di Reti

Ing. Chiara Contoli, PhD

chiara.contoli@unibo.it

*Corso di Laurea Triennale in
Ingegneria e Scienze Informatiche*

Why UDP?

No connection required/established (connection introduces latency)

Simple: no state of the connection needs to be memorized neither in the source nor in the destination

Datagram header are **short:** overhead is reduced

No congestion control: UDP can send data without any control

UDP: User Datagram Protocol (RFC 768)

Transport protocol without “frill”

Provides transport service in a “**best effort**” fashion, UDP datagram can be:

- **Lost**
- Delivered **out of sequence** to the application

Connectionless:

- **No hand-shake** between UDP source and UDP destination
- Each UDP datagram is handled in a **totally independent** manner from other UDP datagrams

UDP Socket Programming

UDP: there is no connection between **client** and **server**

There is no handshake

Each application level message needs to be **contained in a single UDP datagram**

To send a response to the client, server needs to extract from each received datagram the source host **IP address** and the **port number** (of the sender source application)

UDP: data transmitted might not arrive at the destination or might arrive **out of order** compared to the sending sequence

Application point of view: UDP provides an **unreliable** byte transfer service of “groups of bytes” (known as datagrams) between client and server

UDP Socket Programming

UDP: there is no connection between **client** and **server**

There is no handshake

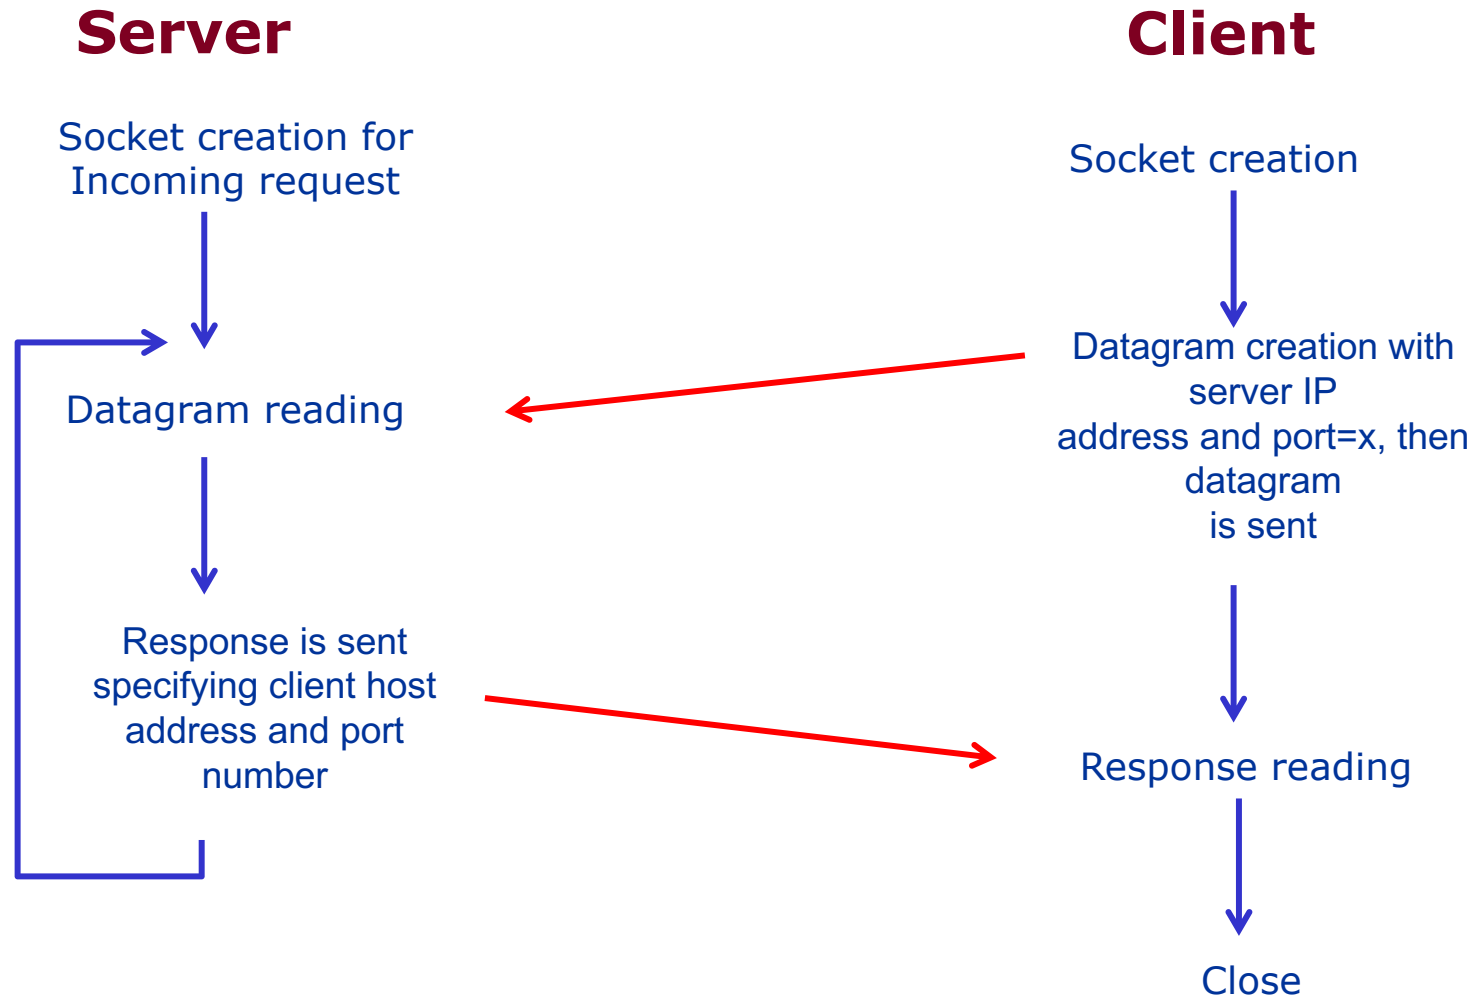
This term is ambiguous, because **datagram** usually is adopted for packets in IP networks, but in the Internet terminology it is adopted pretty often

server needs to extract from each received **datagram** the **IP address** and the **port number** (of the sender source application)

UDP: data transmitted might **not arrive** at the destination or the might arrive **out of order** compared to the sending sequence

Application point of view: UDP provides an **unreliable** byte transfer service of “groups of bytes” (known as datagrams) between client and server

UDP: client – server interaction

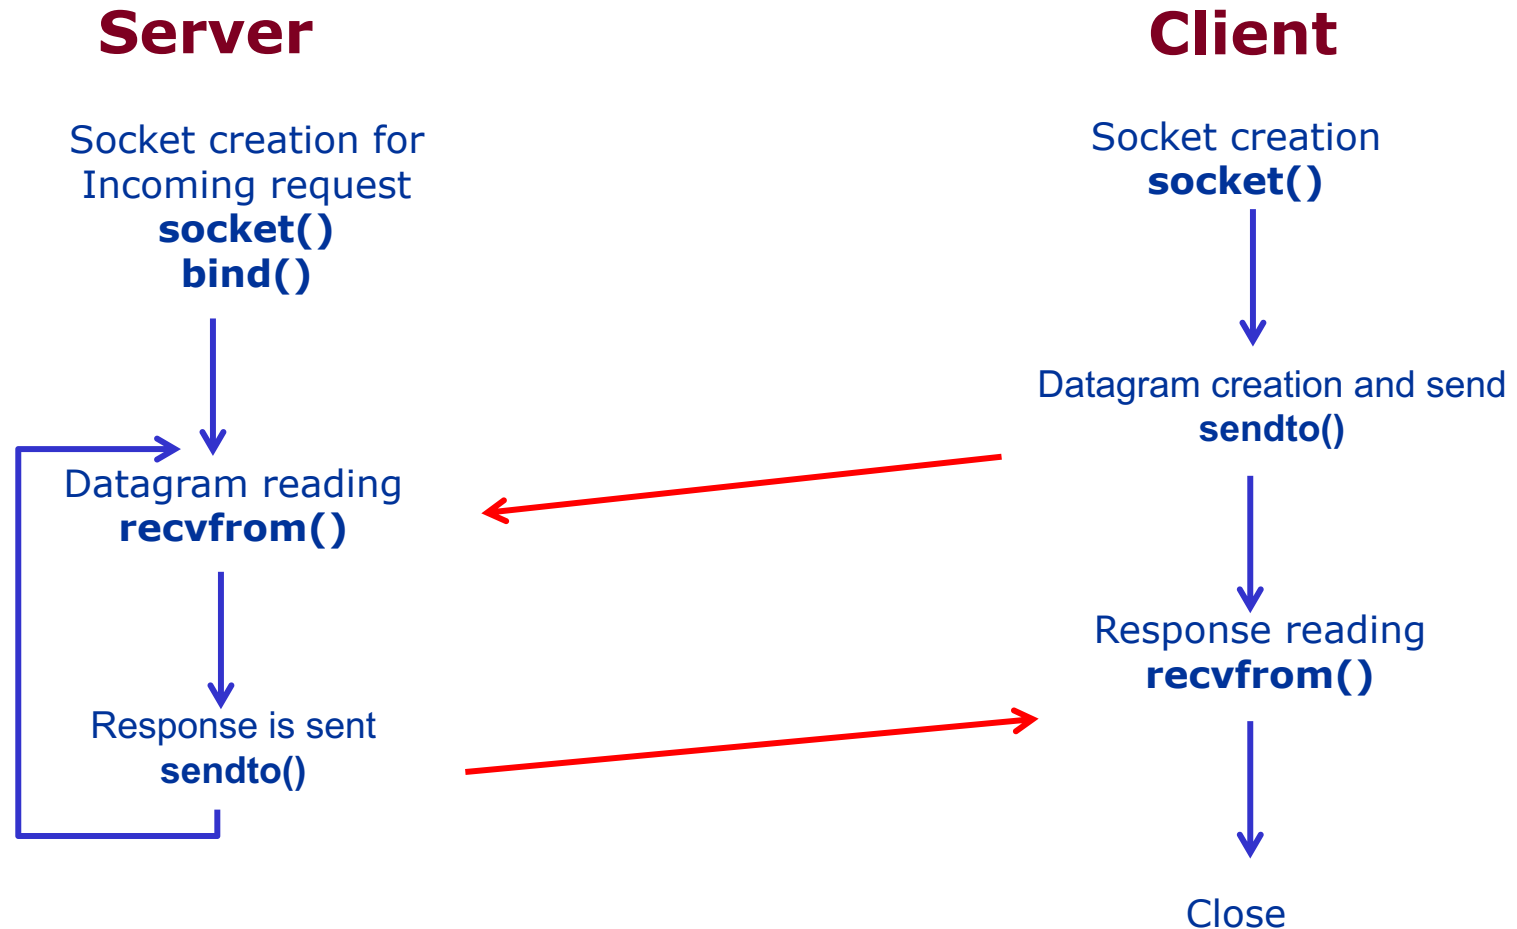


Datagram Socket I/O

From `<sys/socket.h>`

- `ssize_t recvfrom(int sockfd, void *buffer, size_t length, int flags, struct sockaddr *src_addr, socklen_t *addrlen);`
 - Returns number of bytes received, 0 on EOF, -1 on error
 - Parameters: i) *sockfd*: socket file descriptor, ii) *buffer*: store received data, iii) *length*: maximum number of bytes to be read, v) and vi) *src_addr* and *addrlen*: store information about message sender
- `ssize_t sendto(int sockfd, const void *buffer, size_t length, int flags, struct sockaddr *dst_addr, socklen_t addrlen);`
 - Returns number of bytes sent, 0 on EOF, -1 on error
 - Parameters: i) *sockfd*: socket file descriptor, ii) *buffer*: store sent data, iii) *length*: number of bytes to be sent, v) and vi) *src_addr* and *addrlen*: store information about message receiver

UDP: client – server interaction



Example: a simple client / server application

Client:

- User leverage keyboard (standard input) to write a text line
- Client send this text to server

Server:

- Server receive the text line sent by the client
- Convert the text line to upper case letter
- Send modified text to the client

Client:

- Receive modified text line
- Print text on the screen (standard output)

Example: UDP Server

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include "myfunction.h"

#define MAX_BUF_SIZE 1024 // Maximum size of UDP messages
#define SERVER_PORT 9876 // Server port

int main(int argc, char *argv[]){
    struct sockaddr_in server_addr; // struct containing server address
    information
    struct sockaddr_in client_addr; // struct containing client address
    information
    int sfd; // Server socket filed descriptor
    int br; // Bind result
    int i;
    ssize_t byteRecv; // Number of bytes received
    ssize_t byteSent; // Number of bytes to be sent
```

Example: UDP Server

```
socklen_t cli_size;
char receivedData [MAX_BUF_SIZE]; // Data to be received
char sendData [MAX_BUF_SIZE]; // Data to be sent

sfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

if (sfd < 0){
    perror("socket"); // Print error message
    exit(EXIT_FAILURE);
}

// Initialize server address information
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(SERVER_PORT); // Convert to network byte
order
server_addr.sin_addr.s_addr = INADDR_ANY; // Bind to any address

br = bind(sfd, (struct sockaddr *) &server_addr, sizeof(server_addr));
if (br < 0){
    perror("bind"); // Print error message
    exit(EXIT_FAILURE);
}
```

Example: UDP Server

```
cli_size = sizeof(client_addr);
```

```
for(;;){
```

```
    byteRecv = recvfrom(sfd, receivedData, MAX_BUF_SIZE, 0, (struct  
sockaddr *) &client_addr, &cli_size);
```

```
    if(byteRecv == -1){
```

```
        perror("recvfrom");
```

```
        exit(EXIT_FAILURE);
```

```
    }
```

```
    printf("Received data: ");
```

```
    printData(receivedData, byteRecv);
```

```
    if(strncmp(receivedData, "exit", byteRecv) == 0){
```

```
        printf("Command to stop server received\n");
```

```
        break;
```

```
    }
```

```
    convertToUpperCase(receivedData, byteRecv);
```

```
    printf("Response to be sent back to client: ");
```

```
    printData(receivedData, byteRecv);
```


Example: UDP Server

```
byteSent = sendto(sfd, receivedData, byteRecv, 0, (struct sockaddr *)
&client_addr, sizeof(client_addr));

if(byteSent != byteRecv){
    perror("sendto");
    exit(EXIT_FAILURE);
}
} // End of for(;;)

return 0;

}
```

Example: UDP Client

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include "myfunction.h"

#define MAX_BUF_SIZE 1024 // Maximum size of UDP messages
#define SERVER_PORT 9876 // Server port

int main(int argc, char *argv[]){
    struct sockaddr_in server_addr; // struct containing server address
    information
    struct sockaddr_in client_addr; // struct containing client address
    information
    int sfd; // Server socket filed descriptor
    int br; // Bind result
    int i;
    int stop = 0;
```

Example: UDP Client

```
ssize_t byteRecv; // Number of bytes received
ssize_t byteSent; // Number of bytes to be sent

size_t msgLen;
socklen_t serv_size;
char receivedData [MAX_BUF_SIZE]; // Data to be received
char sendData [MAX_BUF_SIZE]; // Data to be sent

sfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

if (sfd < 0){
    perror("socket"); // Print error message
    exit(EXIT_FAILURE);
}

server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(SERVER_PORT);
server_addr.sin_addr.s_addr = inet_addr("127.0.0.1");

serv_size = sizeof(server_addr);
```

Example: UDP Client

```
while(!stop){
    printf("Insert message:\n");
    scanf("%s", sendData);
    printf("String going to be sent to server: %s\n", sendData);

    if(strcmp(sendData, "exit") == 0){
        stop = 1;
    }
    msgLen = countStrLen(sendData);
    byteSent = sendto(sfd, sendData, msgLen, 0, (struct sockaddr *)
&server_addr, sizeof(server_addr));
    printf("Bytes sent to server: %zd\n", byteSent);

    if(!stop){
        byteRecv = recvfrom(sfd, receivedData, MAX_BUF_SIZE, 0, (struct
sockaddr *) &server_addr, &serv_size);
        printf("Received from server: ");
        printData(receivedData, byteRecv);
    }
} // End of while
return 0;
}
```

“myfunction.h”

```
#include <ctype.h>
```

```
size_t countStrLen(char *str){  
    size_t c = 0;
```

```
    while(*str != '\0'){  
        c += 1;  
        str++;  
    }  
    return c;  
}
```

```
void printData(char *str, size_t numBytes){  
    for(int i = 0; i < numBytes; i++){  
        printf("%c", str[i]);  
    }  
    printf("\n");  
}
```

“myfunction.h”

```
void convertToUpperCase(char *str, size_t numBytes){  
    for(int i = 0; i < numBytes; i++){  
        str[i] = toupper(str[i]);  
    }  
}
```

Useful references

Socket API, to cite a few:

- <http://man7.org/linux/man-pages/man3/>
 - <http://man7.org/linux/man-pages/man3/socket.3p.html>
 - <http://man7.org/linux/man-pages/man3/bind.3p.html>
- https://www.gnu.org/software/libc/manual/html_node/Sockets.html
!
- ... and much more on the web
- Kurose, Ross book (Chapter 2, 2.7 and 2.8)

C language reference:

- <http://en.cppreference.com/w/c>